

Constrained Optimization homework (problem 1)

Tommaso Monopoli
s278727

Please refer to the appendix A at the end of this document for all the Matlab codes of this project

1 Problem introduction

Consider the following constrained optimization problem

$$\begin{aligned} \underset{x}{\text{minimize}} \quad & f(x) = \sum_{i=1}^n ix_i^2 \\ \text{subject to} \quad & -5.12 \leq x_i \leq 5.12 \quad (i = 1, \dots, n) \end{aligned} \tag{1}$$

where $n = 10^d$ and $d = 3, 4, 5$.

This problem is strictly convex, in fact:

- $f(x)$ is a strictly convex function (since it is a sum of n strictly convex functions)
- the feasible set X is convex (it is a n -dimensional cube)

therefore, it has a unique solution (in particular, global minimum $f(x^*) = 0$ at $x^* = (0, \dots, 0)$).

One of the many constrained optimization techniques to solve this problem is the projected gradient method (sect. 2).

At each iteration, the projected gradient method computes a descent direction exploiting the gradient of the function $f(x)$, i.e. $\nabla f(x)$. There are two ways of computing the gradient:

- exactly, by computing $\partial f(x)/\partial x_i$ analitically for each i (sect. 3.1)
- numerically, with the so-called finite differences approximation (sect. 3.2)

2 Projected Gradient Method

The projected gradient method is a generalization of the gradient method (aka steepest descent method) to constrained optimization problems.

The main idea of the method is the **projection of a point y onto a set X** :

$$\Pi_X(y) = \arg \min_{x \in X} \|x - y\|^2$$

that is, the point $x \in X$ closest to y .

Computing the projection of a point $y \notin X$ onto X is not an easy task in general (it requires solving a constrained minimization problem), but there are cases in which the structure of X makes the problem easy. For example, if X is a box in $\overline{\mathbb{R}}^n$, i.e. $X = \{x \in \mathbb{R}^n \mid L_i \leq x_i \leq U_i \quad \forall i = 1, \dots, n\}$, then

$$[\Pi_X(y)]_i = \begin{cases} y_i & \text{if } L_i \leq y_i \leq U_i \\ L_i & \text{if } y_i < L_i \\ U_i & \text{if } y_i > U_i \end{cases}$$

In our setting, X is a n -cube with $L_i = -5.12$ and $U_i = +5.12 \quad \forall i = 1, \dots, n$.

The **projected gradient method** implemented for this project is described hereafter.

For $k \geq 0$ (and until a stopping criterion is satisfied):

1. Pick an initial point $x_0 \in X^1$. Set the tolerances for the stopping criteria:
 - $\varepsilon_{\text{grad}}$ (tolerance on the norm of the gradient),
 - $\varepsilon_{\text{diff}}$ (tolerance on the norm of the difference between consecutive iterates),
 - k_{max} (max n. of iterations)
2. Loop until a stopping condition is met:
 - (a) (Steepest Descent Step) pick $-\nabla f(x_k)$ as descent direction and move along it with a steplength $\gamma_k > 0$

$$\hat{x}_k = x_k - \gamma_k \nabla f(x_k)$$

- (b) (Projection Step) compute

$$\bar{x}_k = \Pi_X(\hat{x}_k)$$

- (c) (Update) move along the direction $\bar{x}_k - x_k$ with a steplength $\alpha_k > 0$

$$x_{k+1} = x_k + \alpha_k(\bar{x}_k - x_k)$$

- (d) (Stopping criterion) END if at least one of the following conditions is met:

- $\|\nabla f(x_{k+1})\| < \varepsilon_{\text{grad}}$
- $\|x_{k+1} - x_k\| < \varepsilon_{\text{diff}}$
- $k \geq k_{\text{max}}$

otherwise increment k by 1 and go back to step 2.

The stopping condition $\|x_{k+1} - x_k\| < \varepsilon_{\text{diff}}$ is needed because in a constrained optimization problem a local minimum of f can be attained also in a non-stationary point (i.e. on ∂X).

The choice of the steplengths γ_k and α_k can be done in the following way:

- $\gamma_k = \gamma$ constant $\forall k$

¹if x_0 is given, and it is *outside* X , then initialize $x_0 \leftarrow \Pi_X(x_0)$

- $\alpha_k \in [0, 1]$ chosen with a backtracking line search strategy: with fixed $\rho \in (0, 1)$ and $c_1 \in (0, 1)$

1. start with $\alpha_k^{(0)} = 1$
2. reduce $\alpha_k^{(j)} = \rho \alpha_k^{(j-1)}$ until you satisfy the Armijo condition

$$f\left(x_k + \alpha_k^{(j)}(\bar{x}_k - x_k)\right) \leq f(x_k) + c_1 \alpha_k^{(j)} \nabla f(x_k)^\top (\bar{x}_k - x_k)$$

In all experiments, the maximum number of (inner) iterations of the backtracking line search is bounded to a fixed value k_{\max}^{BT} .

3 Experiment and results

We can clearly see that $f(x)$ is additively separable, in the sense that $f(x) = f_1(x_1) + f_2(x_2) + \dots + f_n(x_n)$, and moreover X is defined by a set of n one-variable inequality constraints, $-5.12 \leq x_i \leq +5.12$, so we could easily solve this optimization problem by dividing it in n smaller constrained optimization problems, which can be solved independently and yield one component x_i^* of the global minimizer x^* .

However, I decided not to solve this problem in such a trivial way, for the sake of generality. Instead, I applied the projected gradient method directly to f .

In all my experiments, I fixed these parameters:

- $c_1 = 10^{-4}$
- $\rho = 0.8$
- $k_{\max}^{\text{BT}} = 100$
- $\varepsilon_{\text{grad}} = 10^{-5}$
- $\varepsilon_{\text{diff}} = 10^{-5}$
- $k_{\max} = 3000$

The initial point x_0 is chosen randomly: each component $x_{0,i}$ is chosen uniformly at random in the interval $[-10, 10]$; therefore, for large n it is very likely that we will need to project x_0 onto X in order for the algorithm to begin. I decided to fix a random seed, for replicability purposes; this means that, given the dimension n , the initial point is always the same.

As for the Projection Step steplength γ , the choice can be done by manual tuning (several values are tried: $\gamma = 1, 0.9, 0.8$).

To visualize how the method behaves, I applied it in the very simple case $n = 2$, with $\gamma = 0.8$. In fig. 1 a 2D plot and 3D plot of f is reported, along with the sequence of iterates leading to the minimum of f .

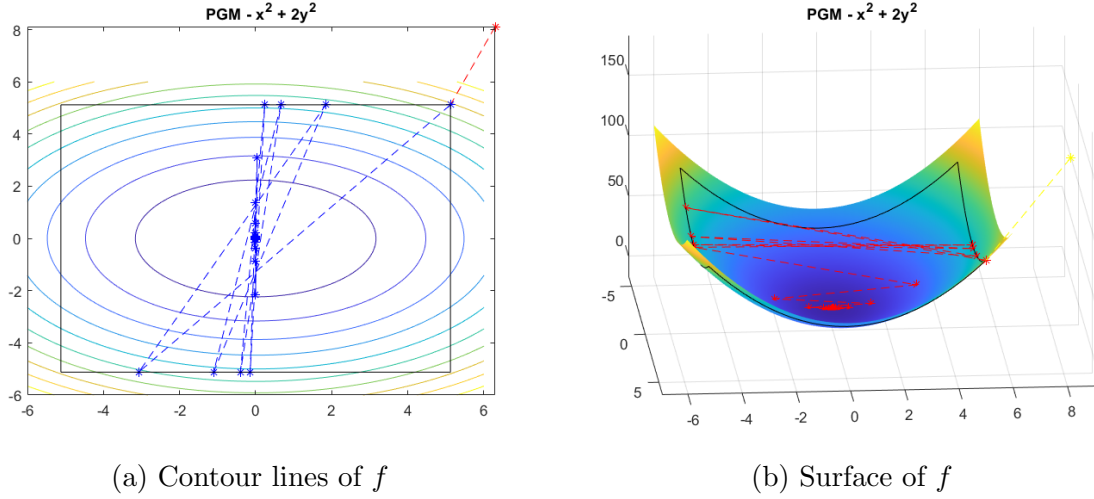


Figure 1: 2D and 3D plots showing how the method behaves when $n = 2$

3.1 Analytic $\nabla f(x_k)$

As mentioned in the introduction, we can apply the projected gradient method to our problem by computing the exact gradient of $f(x)$ analytically

$$[\nabla f(x)]_i = 2i x_i \quad \forall i \quad \implies \quad \nabla f(x) = (2x_1, 4x_2, 6x_3, \dots)$$

I have run the projected gradient method with exact gradient for $n = 10^3, 10^4, 10^5$ and for several values of γ . Results are reported in table 1.

n	γ	#iterations	#projections	$\ \nabla f(x_K)\ $	$\ x_K - x_{K-1}\ $	$f(x_K)$
10^3	1	359	34	0.0097034	9.8045e-06	2.4539e-08
	0.9	68	39	0.0066817	7.724e-06	4.8023e-08
	0.8	362	33	0.0098027	9.9055e-06	2.6611e-08
10^4	1	76	59	0.059506	7.3575e-06	4.7095e-07
	0.9	105	43	0.09199	9.7045e-06	4.0859e-07
	0.8	79	62	0.0723	8.9619e-06	5.0719e-07
10^5	1	85	60	0.75082	8.4509e-06	1.9125e-06
	0.9	104	93	0.64921	8.5169e-06	1.6418e-06
	0.8	86	60	0.7542	8.4832e-06	2.234e-06

Table 1: PGM with exact gradient

It is evident that the speed of convergence of the algorithm is heavily dependent on the choice of γ . However, with a lucky choice of γ , the algorithm converges to the minimum in less than 100 iterations, for all n .

It is interesting to note that the number of (inner) iterations of the backtracking strategy to find α_k is increasing with the number k of outer iterations, and then stabilizes towards the end of the algorithm. In fig. 2 this behaviour is clearly shown (for $n = 10^3$ and $\gamma = 0.9$). This is because the function is steeper far from the minimum point (where early iterates live), and flatter near to the minimum point (and therefore the Armijo condition is satisfied only with smaller α_k , until this need of having a smaller α_k is counterbalanced by $\|\nabla f(x_k)\|$ too becoming very small).

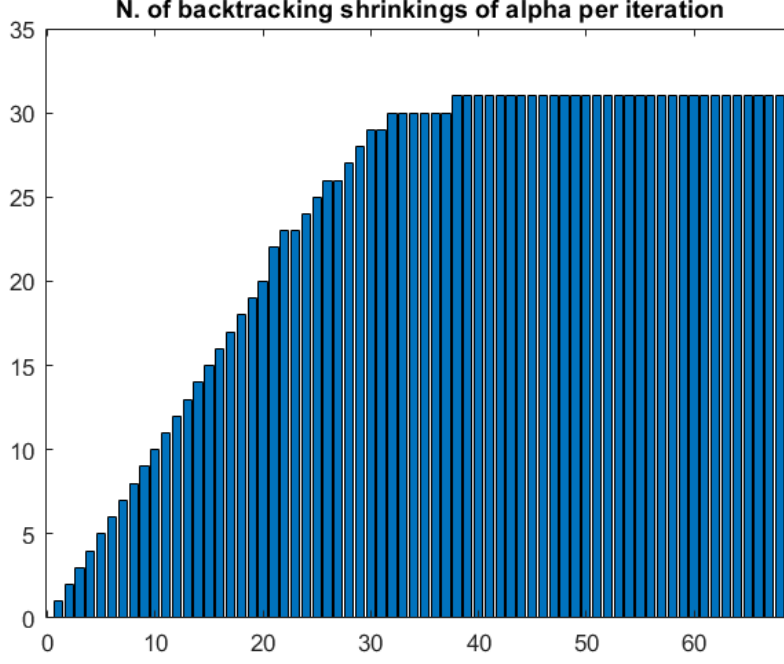


Figure 2: Number of backtracking iterations per outer iteration (for $n = 10^3$ and $\gamma = 0.9$)

3.2 Numerical $\nabla f(x_k)$ (finite differences approximation)

Instead of computing the exact $\nabla f(x)$ analytically, we can approximate it using finite differences approximations of the derivatives.

For our experiment, three versions of the finite differences approximation were employed:

- Forward differences: $\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x+he_i)-f(x)}{h} \quad \forall i = 1, \dots, n$
- Backward differences: $\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x)-f(x-he_i)}{h} \quad \forall i = 1, \dots, n$
- Centered differences: $\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x+he_i)-f(x-he_i)}{2h} \quad \forall i = 1, \dots, n$

The choice of the increment h is of critical importance:

- h too large \implies high analytical error in the estimation of the gradient
- h too small \implies numerical cancellation issues (because the differences $f(x+he_i) - f(x)$, $f(x) - f(x-he_i)$, $f(x+he_i) - f(x-he_i)$ will be very small)

For our problem, we will experiment with an increment $h = 10^{-k}\|x\|$, where $k = 2, 4, 6, 8, 10, 12$.

One issue I encountered, especially when applying the method in a high dimensional space ($n = 10^4, 10^5$), is that computing the gradient with these formulas is very computationally expensive: gradient computation would require n (or $2n$, in the centered differences case) function evaluations for each iteration of the PGM. Therefore, I decided to exploit the particularly simple structure of $f(x)$ to rewrite these formulas in a simpler, more computationally-friendly way.

For the forward difference formula:

$$\begin{aligned}\frac{\partial f(x)}{\partial x_i} &\approx \frac{f(x + he_i) - f(x)}{h} = \frac{\left(\sum_{j=1}^{i-1} jx_j^2 + i(x_i + h)^2 + \sum_{j=i+1}^n jx_j^2\right) - \sum_{j=1}^n jx_j^2}{h} = \\ &= \frac{i(x_i + h)^2 - ix_i^2}{h} = i \frac{2x_i h + h^2}{h} = \mathbf{i}(2\mathbf{x}_i + \mathbf{h})\end{aligned}$$

Similarly, for the backward difference formula:

$$\begin{aligned}\frac{\partial f(x)}{\partial x_i} &\approx \frac{f(x) - f(x - he_i)}{h} = \frac{\sum_{j=1}^n jx_j^2 - \left(\sum_{j=1}^{i-1} jx_j^2 + i(x_i - h)^2 + \sum_{j=i+1}^n jx_j^2\right)}{h} = \\ &= \frac{ix_i^2 - i(x_i - h)^2}{h} = i \frac{2x_i h - h^2}{h} = \mathbf{i}(2\mathbf{x}_i - \mathbf{h})\end{aligned}$$

Finally, for the centered difference formula:

$$\begin{aligned}\frac{\partial f(x)}{\partial x_i} &\approx \frac{f(x + he_i) - f(x - he_i)}{2h} = \\ &= \frac{\left(\sum_{j=1}^{i-1} jx_j^2 + i(x_i + h)^2 + \sum_{j=i+1}^n jx_j^2\right) - \left(\sum_{j=1}^{i-1} jx_j^2 + i(x_i - h)^2 + \sum_{j=i+1}^n jx_j^2\right)}{2h} = \\ &= \frac{i(x_i + h)^2 - i(x_i - h)^2}{2h} = i \frac{4x_i h}{2h} = 2\mathbf{i}\mathbf{x}_i\end{aligned}$$

With these three formulas², the numerical computation of the gradient is almost as fast as the analytic computation. Surprisingly, in the centered differences case the numerical gradient corresponds exactly to the analytic one, whatever h we choose. Therefore, approximating the gradient with the centered differences formula is (theoretically) equivalent to computing it analytically, and so we would obtain the same results already reported in table 1.

Moreover, note that with these new formulas any numerical cancellation problem caused by a very low value of the increment h is avoided.

Nonetheless, for didactic reasons (i.e. for showing off the typical numerical problems arising when choosing h too low), I have employed the usual finite differences formulas whenever computationally feasible, i.e. for $n = 10^3$, and used the new formulas only for the problematic cases $n = 10^4, 10^5$.

²Obviously, this mathematical "trick" is problem-dependent (I exploited the nice structure of $f(x)$) and in no way generalizable to any $f(x)$.

Results are reported in tables 2, 3 and 4 (see page 7). This time, I fixed $\gamma = 0.9$ once and for all, and reported only the n. of iterations and n. of projections for reasons of brevity.

The results are interesting in a number of ways:

- with an adequate choice of h (not too large nor too small), the algorithm converges in a number of iterations very similar to the case with the analytic gradient (despite taking significantly more computational time)
- results in **red** show a very high n. of iterations to converge; this is most likely caused by the value of h being too large, so the numerical gradient at each step is inaccurate w.r.t. the analytic one (so the algorithm is not taking the direction of steepest descent)
- results in **violet** show a very high n. of iterations to converge; since this happens only for $n = 10^3$, this is most likely caused by numerical cancellation issues, which make the numerical gradient again very inaccurate w.r.t. the analytic one
- results in **blue** show a surprisingly low n. of iterations to converge; I verified that in all these cases, the stopping criterion on $\|x_{k+1} - x_k\|$ is met, even though $f(x_k)$ is far from the minimum. This strange behavior might be explained by the fact that the numerical gradient is so inaccurate (due to the very high h) that its opposite is a very "mild" descent direction or not a descent direction at all (in both cases, the backtracking line search yields a very small α_k , therefore $x_{k+1} \approx x_k$)

n	k	#iterations	#projections
10^3	2	2071	38
	4	69	39
	6	68	39
	8	68	39
	10	164	39
	12	2573	39
10^4	2	3000	1140
	4	105	43
	6	105	43
	8	105	43
	10	105	43
	12	105	43
10^5	2	2	3
	4	151	138
	6	98	87
	8	102	91
	10	104	93
	12	104	93

Table 2: PGM with numerical gradient (forward differences approximation)

n	k	#iterations	#projections
10^3	2	1987	43
	4	68	39
	6	68	39
	8	68	39
	10	433	39
	12	2117	39
10^4	2	53	51
	4	106	43
	6	105	43
	8	105	43
	10	105	43
	12	105	43
10^5	2	2	3
	4	167	152
	6	104	93
	8	104	93
	10	104	93
	12	104	93

Table 3: PGM with numerical gradient (backward differences approximation)

n	k	#iterations	#projections
10^3	2	68	39
	4	68	39
	6	68	39
	8	68	39
	10	68	39
	12	624	39
10^4	2	105	43
	4	105	43
	6	105	43
	8	105	43
	10	105	43
	12	105	43
10^5	2	104	93
	4	104	93
	6	104	93
	8	104	93
	10	104	93
	12	104	93

Table 4: PGM with numerical gradient (centered differences approximation)

A Matlab code

In this section is reported the Matlab code produced for this project

A.1 pgm_test.m

```
1 clear
2 close all
3 clc
4
5 %% DEFINE THE PROBLEM
6
7 n = 10^3; % n. of variables
8 f = @(x) [1:n] * (x.^2); % f. handle for the objective function
9 gradf = @(x) [1:n]' .* (2*x); % f. handle for the gradient of the ...
    objective function
10 box_mins = -5.12 * ones(n,1); % upper boundaries of the feasible box
11 box_maxs = 5.12 * ones(n,1); % lower boundaries of the feasible box
12
13
14
15 %% DEFINE PARAMETERS FOR THE PROJECTED GRADIENT METHOD
16
17 rng(0); % set seed (for replicability)
18 x0 = unifrnd(-10,10,n,1);
19
20 % for the projection step
21 gamma = 0.9; % 0.8, 0.9, 1
22 Pi_X = @(x) max(min(x, box_maxs), box_mins); % f. handle for the ...
    box projection function
23
24 % For the purpose of comparing the projected gradient method and the
25 % unconstrained gradient method, here is a function handle for an ...
    identity
26 % projection function
27 %Pi_X = @(x) x;
28
29 % for the backtracking line search
30 c1 = 1e-4;
31 rho = 0.8;
32 btmax = 100;
33
34 % for the stopping criterion
35 tolgrad = 1e-5; % iterate until ||gradf(x(k))|| < tolgrad
36 tolx = 1e-5; % iterate until ||x(k+1)-x(k)|| < tolx
37 kmax = 3000;
38
39 % Choose how to compute the gradient
40 % - 'forw' --> approximate grad(f) with the forward difference ...
    approximation
41 % - 'backw' --> approximate grad(f) with the backward difference ...
    approximation
42 % - 'centr' --> approximate grad(f) with the centered difference ...
    approximation
43 % - 'exact' (or just every other string) --> use the exact grad(f)
```

```

44 FDgrad = 'exact';
45 k_FDgrad = 2; % the exponent which appears in the expression of the ...
    increment h=10^(-k) * ||x||
46
47 verbose = true;      % if true, print the current iteration every 10
48
49
50
51 %% RUN THE PGM
52
53 disp('**** CONSTR. STEEPEST DESCENT: START ****')
54 tic      % start timer
55
56 [xk_final, fk_final, gradfk_norm_final, deltaxk_norm_final, k_final, ...
57  xseq_final, btseq_final, projection_count] = ...
58  projected_gradient_method(x0, f, gradf, kmax, tolgrad, c1, rho, ...
59  btmax, gamma, tolX, Pi_X, verbose, FDgrad, k_FDgrad);
60
61 disp('**** CONSTR. STEEPEST DESCENT: FINISHED ****')
62 disp('')
63
64 disp('**** CONSTR. STEEPEST DESCENT: RESULTS ****')
65 disp('*****')
66 toc      % stop timer
67 if n==2
68     disp(['xk: ', mat2str(xk_final), ' (actual minimizer: vector of ...
        0s);'])
69 end
70 disp(['norm(xk): ', mat2str(norm(xk_final)), ' (actual norm of the ...
    minimizer: 0; NB: this is also the norm of the error, ||xk-x*||);'])
71 disp(['f(xk): ', num2str(fk_final), ' (actual min. value: 0);'])
72 disp(['N. of iterations: ', num2str(k_final), '/', num2str(kmax), ';'])
73 disp(['N. of projections: ', num2str(projection_count), ';'])
74 disp(['gradient norm: ', num2str(gradfk_norm_final), ';'])
75 disp(['length of last step: ', num2str(deltaxk_norm_final), ';'])
76 disp('*****')
77
78
79
80 %% PLOTS
81
82 % Barplot of btseq
83 fig_bt_iters = figure();
84 bar(btseq_final)
85 title('PGM - x^2 + 2y^2')
86
87 if n==2      % in this very simple case, we can visualize the ...
    objective function in 3D
88
89     % Creation of the data to plot the domain boundaries
90     t = linspace(0, 1, 25);
91     dom_xy_1 = box_mins + t .* ([box_mins(1); box_maxs(2)] - box_mins);
92     dom_xy_2 = [box_mins(1); box_maxs(2)] + t .* (box_maxs - ...
        [box_mins(1); box_maxs(2)]);
93     dom_xy_3 = box_maxs + t .* ([box_maxs(2); box_mins(1)] - box_maxs);
94     dom_xy_4 = [box_maxs(2); box_mins(1)] + t .* (box_mins - ...
        [box_maxs(2); box_mins(1)]);
95

```

```

96     dom_xy = [dom_xy_1, dom_xy_2, dom_xy_3, dom_xy_4];
97     f_z = f(dom_xy);
98
99     % Projection of the starting point
100    Pi_X_x0 = Pi_X(x0);
101
102    % Creation of the meshgrid for the contour-plot
103    [X, Y] = meshgrid(linspace(-6, 6, 500), linspace(-6, 6, 500));
104
105    % Computation of the values of f for each point of the mesh
106    Z = X.^2 + 2*Y.^2;
107
108    % Simple Plot
109    fig_contour = figure();
110    % Contour plot with curve levels for each point in xseq
111    [C1, ~] = contour(X, Y, Z);
112    hold on
113    % plot of the points in xseq
114    plot([x0(1), Pi_X_x0(1)], [x0(2), Pi_X_x0(2)], 'r--*')
115    plot([Pi_X_x0(1) xseq_final(1, :)], [Pi_X_x0(2) xseq_final(2, ...
116        :)], 'b--*')
116    plot(dom_xy(1, :), dom_xy(2, :), 'k')
117    hold off
118    title('PGM - x^2 + 2y^2')
119
120    % Much more interesting plot
121    fig_surface = figure();
122    surf(X, Y, Z, 'EdgeColor', 'none')
123    hold on
124    plot3([x0(1) Pi_X_x0(1)], [x0(2) Pi_X_x0(2)], [f(x0), ...
125        f(Pi_X_x0)], 'y--*')
125    plot3([Pi_X_x0(1) xseq_final(1, :)], [Pi_X_x0(2) xseq_final(2, ...
126        :)], [f(Pi_X_x0), f(xseq_final)], 'r--*')
126    plot3(dom_xy(1, :), dom_xy(2, :), f_z, 'k')
127    hold off
128    title('PGM - x^2 + 2y^2')
129
130 end

```

A.2 projected_gradient_method.m

```

1 function [xk, fk, gradfk_norm, deltaxk_norm, k, xseq, btseq, ...
2     projection_count] = ...
3     projected_gradient_method(x0, f, gradf, ...
4         kmax, tolgrad, c1, rho, btmax, gamma, tolx, Pi_X, verbose, ...
5         FDgrad, kFDgrad)
6
7 %
8 % Projected gradient method (steepest descent) for constrained ...
9 % optimization.
10
11 %
12 % INPUTS:
13 % x0 = n-dimensional column vector;
14 % f = function handle that describes a function  $\mathbb{R}^n \rightarrow \mathbb{R}$ ;
15 % gradf = function handle that describes the gradient of f;
16 % kmax = maximum number of iterations permitted;
17 % tolgrad = value used as stopping criterion w.r.t. the norm of the

```

```

13 %             gradient;
14 % c1 = the factor of the Armijo condition that must be a scalar in ...
    (0,1);
15 % rho = fixed factor, less than 1, used for reducing alpha0;
16 % btmax = maximum number of steps for updating alpha during the
17 %         backtracking strategy.
18 % gamma = the initial factor that multiplies the descent direction ...
    at each
19 %         iteration;
20 % tolX = value used as stopping criterion w.r.t. the norm of the
21 %         steps (xnew - xk); infact, in constrained optimization the ...
    global
22 %         minimizer could also be a non-stationary point, so the stopping
23 %         criterion based on a tolerance on the norm of the gradient ...
    is not
24 %         sufficient attained in a non-stationary point
25 % Pi_X = projection function
26 % verbose = if true, print the current iteration every 10
27 % FDgrad = choose how to compute the gradient (exactly or approximately)
28 % k_FDgrad = the exponent which appears in the expression of the ...
    increment
29 %          $h=10^{(-k\_FDgrad)} * ||x||$ 
30 %
31 % OUTPUTS:
32 % xk = the last x computed by the function;
33 % fk = the value f(xk);
34 % gradfk_norm = value of the norm of gradf(xk)
35 % deltaxk_norm = length of the last step of the sequence
36 % k = index of the last iteration performed
37 % xseq = n-by-k matrix where the columns are the xk computed during the
38 % iterations
39 % btseq = 1-by-k vector where elements are the number of backtracking
40 % iterations at each optimization step.
41
42 switch FDgrad
43     case 'forw'
44         % overwrite gradf with a f. handle that uses the forward ...
            difference
45         % approximation
46         h = @(x) (10^(-k_FDgrad))*norm(x);
47         n = size(x0,1);
48
49         if n < 10^4
50             % general-purpose forw. diff. approx.
51             gradf = @(x) (f(x + h(x)*eye(n)) - f(x))' / h(x);
52         else
53             % problem-dependent forw. diff. approx.
54             gradf = @(x) [1:n]' .* (2*x + h(x));
55         end
56
57     case 'backw'
58         % overwrite gradf with a f. handle that uses the backward ...
            difference
59         % approximation
60         h = @(x) (10^(-k_FDgrad))*norm(x);
61         n = size(x0,1);
62
63         if n < 10^4

```

```

64         % general-purpose backw. diff. approx.
65         gradf = @(x) (f(x) - f(x - h(x)*eye(n)))' / h(x);
66     else
67         % problem-dependent backw. diff. approx.
68         gradf = @(x) [1:n]' .* (2*x - h(x));
69     end
70
71     case 'centr'
72         % overwrite gradf with a f. handle that uses the centered ...
73         % approximation
74         h = @(x) (10^(-k_FDgrad))*norm(x);
75         n = size(x0,1);
76
77         if n < 10^4
78             % general-purpose centr. diff. approx.
79             gradf = @(x) (f(x + h(x)*eye(n)) - f(x - h(x)*eye(n)))' ...
80                 / (2*h(x));
81         else
82             % problem-dependent centr. diff. approx.
83             gradf = @(x) [1:n]' .* (2*x); % exact gradient!
84         end
85     otherwise
86         % we use the input function handle gradf
87     end
88
89     % Function handle for the armijo condition
90     farmijo = @(fk, alpha, gradfk, pk) fk + c1 * alpha * gradfk' * pk;
91
92     % Initializations
93     xseq = zeros(length(x0), kmax);
94     btseq = zeros(1, kmax);
95     projection_count = 0;
96
97     xk = Pi_X(x0); % Project the starting point if outside the constraints
98     if ~isequal(xk, x0)
99         projection_count = projection_count + 1;
100 end
101 fk = f(xk);
102
103 % compute the gradient (exactly or approximately, depending on FDgrad)
104 gradfk = gradf(xk);
105
106 k = 0;
107 gradfk_norm = norm(gradfk);
108 deltaxk_norm = tol_x + 1; % this is to ensure that at least a ...
109     first iteration is performed
110
111 if verbose
112     disp(['iteration: ', num2str(k)])
113     disp(['norm(xk): ', num2str(norm(xk))])
114     disp(['gradient norm: ', num2str(gradfk_norm)])
115     if strcmp(FDgrad, 'forw') || strcmp(FDgrad, 'backw') || ...
116         strcmp(FDgrad, 'centr')
117         disp(['h(xk): ', num2str(h(xk))])
118     end
119 end

```

```

118
119 while k < kmax && gradfk_norm >= tolgrad && deltaxk_norm >= tolx
120     % Compute the descent direction (exactly or approximately, ...
121     % depending on FDgrad)
122     pk = -gradf(xk);
123     % Take a step in the descent direction and project the resulting ...
124     % vector
125     % onto the feasible set X
126     xhatk = xk + gamma * pk
127     xbark = Pi_X(xhatk);
128     if ~isequal(xbark, xhatk)
129         projection_count = projection_count + 1;
130     end
131
132     % Backtracking line search
133
134     % Reset the value of alpha
135     alpha = 1;
136
137     % Compute the candidate new xk
138     pik = xbark - xk;
139     xnew = xk + alpha * pik;
140
141     % Compute the value of f in the candidate new xk
142     fnew = f(xnew);
143
144     bt = 0;      % counter of backtracking iterations
145
146     % 2nd condition is the Armijo (w.r.t. pik) condition not satisfied
147     while bt < btmax && fnew > farmijo(fk, alpha, gradfk, pik)
148         % Reduce the value of alpha
149         alpha = rho * alpha;
150         % Update xnew and fnew w.r.t. the reduced alpha
151         xnew = xk + alpha * pik;
152         fnew = f(xnew);
153
154         % Increase the counter by one
155         bt = bt + 1;
156
157     end
158
159     % Update xk, fk, gradfk_norm, deltaxk_norm
160     deltaxk_norm = norm(xnew - xk);
161     xk = xnew;
162     fk = fnew;
163     gradfk = gradf(xk);
164     gradfk_norm = norm(gradfk);
165
166     % Increase the step by one
167     k = k + 1;
168
169     % Store current xk in xseq
170     xseq(:, k) = xk;
171     % Store bt iterations in btseq
172     btseq(k) = bt;
173

```

```

174     if verbose && mod(k,10) == 0
175         disp(['iteration: ', num2str(k)])
176         disp(['norm(xk): ', num2str(norm(xk))])
177         disp(['gradient norm: ', num2str(gradfk_norm)])
178         if strcmp(FDgrad, 'forw') || strcmp(FDgrad, 'backw') || ...
            strcmp(FDgrad, 'centr')
179             disp(['h(xk): ', num2str(h(xk))])
180         end
181     end
182 end
183
184 % "Cut" xseq and btseq to the correct size
185 xseq = xseq(:, 1:k);
186 btseq = btseq(1:k);
187
188 end

```